

```
In [1]: %matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.metrics import r2_score
import statsmodels.api as sm
from sklearn.linear_model import LinearRegression
import time
```

Univariate Linear Regression

```
In [2]: data=pd.read_csv("D:/Sem 5/ML/exiddata1.txt", header=None)
data.head()
```

```
Out[2]:
```

	0	1
0	6.1101	17.5920
1	5.5277	9.1302
2	8.5186	13.6620
3	7.0032	11.8540
4	5.8598	6.8233

```
In [3]: data.describe()
```

```
Out[3]:
```

	0	1
count	97.000000	97.000000
mean	8.159800	5.839135
std	3.869884	5.510262
min	5.026900	-2.680700
25%	5.707700	1.986900
50%	6.589400	4.562300
75%	8.578100	7.046700
max	22.203000	24.147000

```
In [4]: data.columns = ['Population', 'Profit'] #assigning column names
```

```
In [5]: data.isnull().sum()
```

```
Out[5]: Population    0
Profit              0
dtype: int64
```

```
In [6]: plt.scatter(data['Population'], data['Profit'])
plt.xticks(np.arange(5,30,step=5))
plt.yticks(np.arange(-5,30,step=5))
plt.xlabel('Population (in 10,000s)')
plt.ylabel('Profit (in 10,000s)')
plt.title('Profit Vs Population')
```

```
Out[6]: Text(0.5, 1.0, 'Profit vs Population')
```

Cost Function J(θ)

```
In [7]: def computeCost(X,y,theta):
"""
    Take in a numpy array X,y,theta and get cost function using theta as
    parameter in a linear regression model
    """
    m = len(y)
    predictions = X.dot(theta)
    square_err = (predictions - y)**2
    return 1/(m)*np.sum(square_err)
```

```
In [8]: data['x0'] = 1
```

```
In [9]: data_val = data.values
m = len(data_val[:-1])
X = data[['x0','Population']].iloc[:-1].values
y = data['Profit'].iloc[:-1].values.reshape(m,1)
theta = np.zeros((2,1))

m, X.shape, y.shape, theta.shape
```

```
Out[9]: (96, (96, 2), (96, 1), (2, 1))
```

```
In [10]: # h(θ) = x0θ0 + x1θ1 .... (x0 = 1)
```

```
In [11]: computeCost(X,y,theta)
```

```
Out[11]: 64.80968355754062
```

Gradient Descent

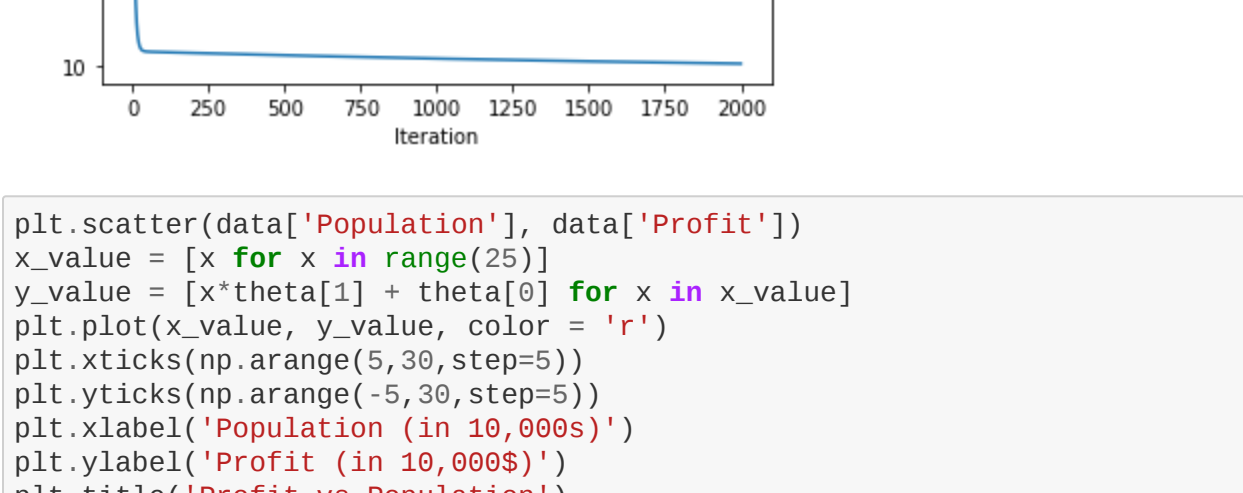
```
In [12]: def gradientDescent(X,y,theta,alpha,num_iters):
"""
    Take numpy array for X,y,theta and update theta for every iteration
    of gradient steps
    Return theta and the list of cost of theta during each iteration
    """
    m = len(y)
    J_history = []
    for i in range(num_iters):
        predictions = X.dot(theta)
        error = np.dot(X.transpose(), (predictions-y))
        descent = alpha * 1/m * error
        theta -=descent
        J_history.append(computeCost(X,y,theta))

    return theta, J_history
```

```
In [13]: theta, J_history = gradientDescent(X,y,theta,0.001, 2000)
```

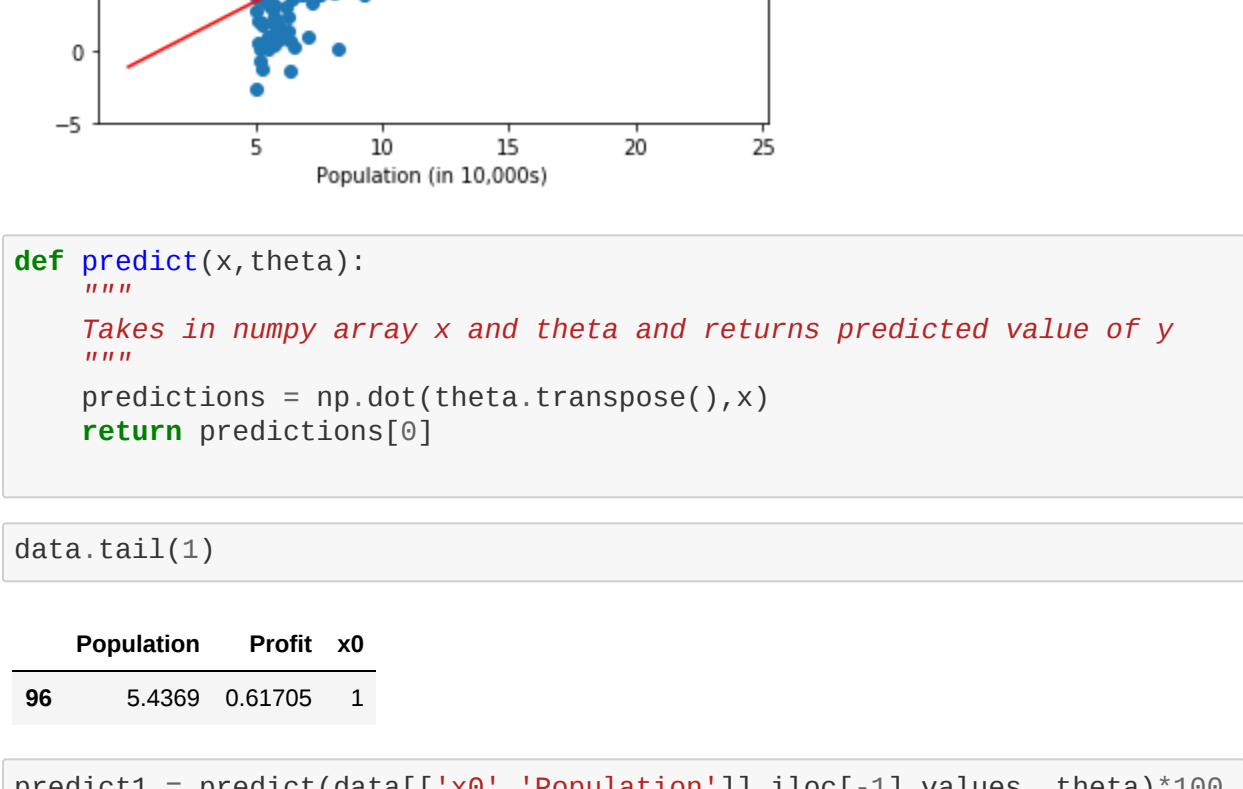
```
In [14]: print(f'h(x) = {str(round(theta[0,0],2))} + {str(round(theta[1,0],2))}x
1')
h(x) = -1.11 + 0.92x1
```

```
In [15]: from mpl_toolkits.mplot3d import Axes3D
#Generating values for theta0, theta1 and the resulting cost value
theta0_vals=np.linspace(-10,10,100)
theta1_vals=np.linspace(-1,4,100)
J_vals=np.zeros((len(theta0_vals),len(theta1_vals)))
for i in range(len(theta0_vals)):
    for j in range(len(theta1_vals)):
        t=np.array([theta0_vals[i],theta1_vals[j]])
        J_vals[i,j]=computeCost(X,y,t)
#Generating the surface plot
fig=plt.figure()
ax=fig.add_subplot(111, projection='3d')
surf=ax.plot_surface(theta0_vals,theta1_vals,J_vals,cmap='coolwarm')
fig.colorbar(surf, shrink=0.5, aspect=5)
ax.set_xlabel("$\theta_0$")
ax.set_ylabel("$\theta_1$")
ax.set_zlabel("$J(\theta_0,\theta_1)$")
#rotate for better angle
ax.view_init(30,120)
```



```
In [16]: plt.plot(J_history)
plt.xlabel('Iteration')
plt.ylabel('$J(\theta_0,\theta_1)$')
plt.title('Cost function using Gradient Descent')
```

```
Out[16]: Text(0.5, 1.0, 'Cost function using Gradient Descent')
```



```
In [17]: plt.scatter(data['Population'], data['Profit'])
x_value = [x for x in range(25)]
y_value = [x*theta[1] + theta[0] for x in x_value]
plt.plot(x_value, y_value, color = 'r')
plt.xticks(np.arange(5,30,step=5))
plt.yticks(np.arange(-5,30,step=5))
plt.xlabel('Population (in 10,000s)')
plt.ylabel('Profit (in 10,000s)')
plt.title('Profit vs Population')
```

```
Out[17]: Text(0.5, 1.0, 'Profit vs Population')
```



```
In [18]: def predict(x,theta):
"""
    Takes in numpy array x and theta and returns predicted value of y
    """
    predictions = np.dot(theta.transpose(),x)
    return predictions[0]
```

```
In [19]: data.tail(1)
```

```
Out[19]:
```

	Population	Profit	x0
96	5.4369	0.61705	1

```
In [20]: predict = predict(data[['x0','Population']].iloc[-1].values, theta)*100
80
print(f'For a population of 6170 the predicted profit is ${predict}')
For a population of 6170 the predicted profit is $38686.246103378166
```

TODO: Configure code for multivariate linear regression

```
In [21]: hw = pd.read_csv("D:/Sem 5/ML/exiddata2.txt", header=None)
hw.head()
```

```
Out[21]:
```

	0	1	2
0	2104	3	399900
1	1600	3	329900
2	2400	3	369000
3	1416	2	232000
4	3000	4	539900

```
In [22]: hw.columns = ["Size","Num_bedrooms","Price"]
hw.head(1)
```

```
Out[22]:
```

	Size	Num_bedrooms	Price
0	2104	3	399900

```
In [23]: hw.describe()
```

```
Out[23]:
```

	Size	Num_bedrooms	Price
count	47.000000	47.000000	47.000000
mean	2000.680851	3.170213	340412.659574
std	794.702354	0.760882	125039.899586
min	852.000000	1.000000	169900.000000
25%	1432.000000	3.000000	249900.000000
50%	1688.000000	3.000000	299900.000000
75%	2269.000000	4.000000	384450.000000
max	4478.000000	5.000000	699900.000000

```
In [24]: hw.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 47 entries, 0 to 46
Data columns (total 3 columns):
 # Column Non-Null Count Dtype
---
 0 Size 47 non-null int64
 1 Num_bedrooms 47 non-null int64
 2 Price 47 non-null int64
dtypes: int64(3)
memory usage: 1.2 KB
```

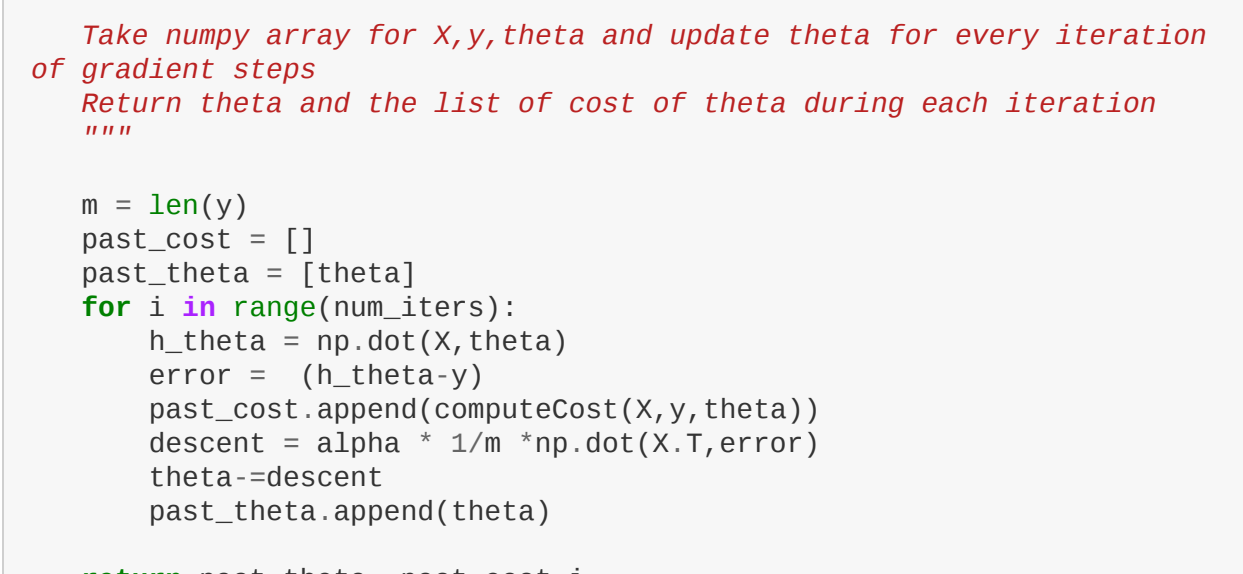
```
In [25]: hw.isnull().sum()
```

```
Out[25]: Size 0
Num_bedrooms 0
Price 0
dtype: int64
```

Price vs Size of house and Price vs No. of bedrooms

```
In [26]: plt.scatter(hw['Size'], hw['Price'])
plt.xticks(np.arange(1000,500,step=500))
plt.yticks(np.arange(100000,800000,step=50000))
plt.xlabel('Size')
plt.ylabel('Price')
plt.title('Price vs Size')
```

```
Out[26]: Text(0.5, 1.0, 'Price vs Size')
```



```
In [27]: plt.scatter(hw['Num_bedrooms'], hw['Price'])
plt.xticks(np.arange(0,6,step=1))
plt.yticks(np.arange(-5,30,step=5))
plt.xlabel('Num_bedrooms')
plt.ylabel('Price')
plt.title('Price vs Num_bedrooms')
```

```
Out[27]: Text(0.5, 1.0, 'Price vs Num_bedrooms')
```



Scaling the column "Size"

```
In [28]: from sklearn.preprocessing import MinMaxScaler
col = ["Size","Price"]
scaler = MinMaxScaler()

hw[col] = pd.DataFrame(scaler.fit_transform(hw[col]),columns = hw[col].columns)
```

Defining X and y

```
In [29]: X = np.array(hw.drop("Price",axis=1))[:-1])
y = np.array(hw["Price"][:-1])
X.shape,y.shape
```

```
Out[29]: ((46, 2), (46,))
```

```
In [30]: y = y.reshape(y.shape[0],1)
X = np.c_[np.ones(X.shape[0]),X]
X.shape, y.shape
```

```
Out[30]: ((46, 3), (46, 1))
```

Defining functions

```
In [31]: def computeCost(X,y,theta):
"""
    Take in a numpy array X,y,theta and get cost function using theta as
    parameter in a linear regression model
    """
    m = y.size
    h_theta = np.dot(X,theta) #Predictions
    error = (h_theta - y)
    cost = (1/(2*m))*np.dot(error.T,error)
    return cost
```

```
In [32]: def gradientDescent(X,y,theta,alpha,num_iters):
"""
    Take numpy array for X,y,theta and update theta for every iteration
    of gradient steps
    Return theta and the list of cost of theta during each iteration
    """
    m = len(y)
    past_cost = []
    theta = [theta]
    for i in range(num_iters):
        h_theta = np.dot(X,theta)
        error = (h_theta - y)
        past_cost.append(computeCost(X,y,theta))
        descent = alpha * 1/m * np.dot(X.T,error)
        theta -=descent
        past_theta.append(theta)

    return past_theta, past_cost,i
```

```
In [45]: def predict(x,theta):
"""
    Takes in numpy array x and theta and returns predicted value of y
    """
    predictions = np.dot(theta.T,x)
    return predictions[0]
```

Declaring θ and α and number of iterations

```
In [34]: np.random.seed(1111)
theta = np.random.rano(X.shape[1],1)
alpha = 0.01
num_iters = 5000
```

Performing regression

```
In [35]: import time
start = time.time()
past_theta,past_cost,stop = gradientDescent(X,y,theta,alpha,num_iters)
timeTaken = time.time() -start
```

```
In [36]: best_theta = past_theta[-1]
best_cost = past_cost[-1]
```

```
In [37]: print(f'The model performed {stop} iterations out of {num_iters}.')
print(f'Best Theta: {best_theta}')
print(f'Best cost: {best_cost}')
```

The model performed 4999 iterations out of 5000.
Best Theta: [[0.06027362]
[0.93272880]
[-0.01125156]]
Best cost: [[0.00743669]]

```
In [53]: print(f'h(theta) = {str(round(best_theta[0,0],5))} + {str(round(best_theta[0,0],5))}x1 + {str(round(best_theta[0,0],5))}x2')
print(f'Time taken: {timeTaken}')
print(f'Accuracy: {round(r2_score(y,predict(X,best_theta,4)))*100}%')
```

h(theta) = 0.06027 + 0.06027x1 + 0.86027x2
Time taken: 0.05385303497314453

```
ValueError                                Traceback (most recent call la
st)
<ipython-input-53-fe1c492f4b65> in <module>
      1 print(f'h(theta) = {str(round(best_theta[0,0],5))} + {str(round(
best_theta[0,0],5))}x1 + {str(round(best_theta[0,0],5))}x2')
----> 3 print(f'Time taken: {timeTaken}')
----> 3 print(f'Accuracy: {round(r2_score(y,predict(X,best_theta,4)))*100}%')
```

```
<ipython-input-45-e8c352d812f8> in predict(x, theta)
      3 Takes in numpy array x and theta and returns predicted value
of y
      4 """
----> 5 predictions = np.dot(theta.T,x)
      6 return predictions[0]
```

```
<_array_function__ internals> in dot(*args, **kwargs)
ValueError: shapes (1,3) and (46,3) not aligned: 3 (dim 1) != 46 (dim 0)
```

Plotting cost

```
In [40]: cost = np.asarray(past_cost).reshape((len(past_cost),1))
cost.shape
```

```
Out[40]: (5000, 1)
```

```
In [ ]:
```

```
In [41]: plt.plot(cost)
plt.show
```

```
Out[41]: <function matplotlib.pyplot.show(*args, **kw)>
```


Predicting price for test data

```
In [42]: hw.tail(1)
```

```
Out[42]:
```

	Size	Num_bedrooms	Price
46	0.096001	3	0.131321

```
In [47]: from sklearn.metrics import r2_score
prediction = predict(X[-1],best_theta)*1000000
print(f'For house of size 968 sq.ft the predicted price is :${str(predic
tion)}')
```

For house of size 968 sq.ft the predicted price is :\$273328.3060931557

```
ValueError                                Traceback (most recent call la
st)
<ipython-input-47-1456dc3b3729> in <module>
      2 prediction = predict(X[-1],best_theta)*1000000
      3 print(f'For house of size 968 sq.ft the predicted price is :${st
r(prediction)}')
```

```
<ipython-input-45-e8c352d812f8> in predict(x, theta)
      3 Takes in numpy array x and theta and returns predicted value
of y
      4 """
----> 5 predictions = np.dot(theta.T,x)
      6 return predictions[0]
```

```
<_array_function__ internals> in dot(*args, **kwargs)
ValueError: shapes (1,3) and (46,3) not aligned: 3 (dim 1) != 46 (dim 0)
```

Comparing the parameters

```
In [44]: print(f'Parameters from StatsModels: {sm.OLS(y,X).fit().params()}'
print(f'Parameters from ScikitLearn: {LinearRegression().fit(X,y).coef
...}')
print(f'Parameters from GradientDescent: {best_theta.reshape(3,)}')
```

```
-----
TypeError                                Traceback (most recent call la
st)
<ipython-input-44-09eb6f790d7c> in <module>
----> 1 print(f'Parameters from StatsModels: {sm.OLS(y,X).fit().param
s()}'
      2 print(f'Parameters from ScikitLearn: {LinearRegression().fit(X,
y).coef[3]}'
      3 print(f'Parameters from GradientDescent: {best_theta.reshape
((3,))}')
```

TypeError: 'numpy.ndarray' object is not callable

```
In [ ]:
```